

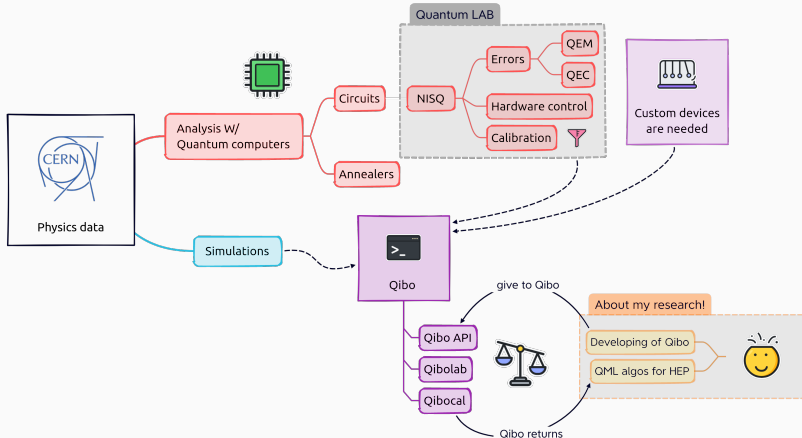
Doing full stack QML using qibo

Matteo Robbiati **feat.** Alejandro Sopena

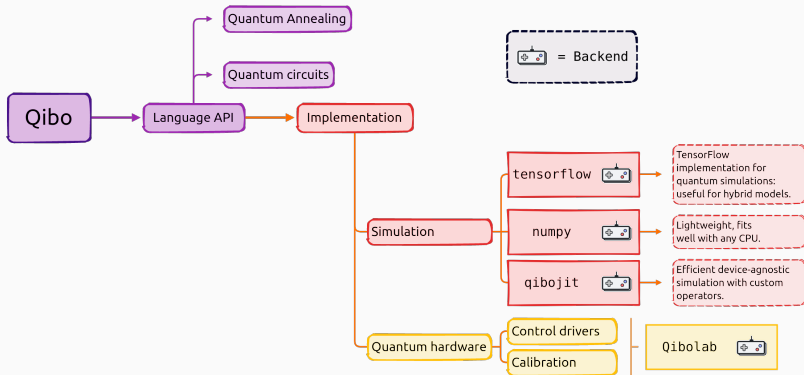
24 May 2023



Working in the NISQ era



The qibo ecosystem



Two snapshots

- ➡ Consider two Hamiltonians H_0 and H_1 , whose ground states are $|g_0\rangle$ and $|g_1\rangle$.

Adiabatic Evolution

- ➔ Consider two Hamiltonians H_0 and H_1 , whose ground states are $|g_0\rangle$ and $|g_1\rangle$.
- ➔ We call Adiabatic Evolution the process represented by:

Adiabatic Evolution

- ➡ Consider two Hamiltonians H_0 and H_1 , whose ground states are $|g_0\rangle$ and $|g_1\rangle$.
- ➡ We call Adiabatic Evolution the process represented by:

$$H_{ad}(\tau; \theta) = [1 - s(\tau; \theta)]H_0 + s(\tau; \theta)H_1. \quad (1)$$

Adiabatic Evolution

- ➡ Consider two Hamiltonians H_0 and H_1 , whose ground states are $|g_0\rangle$ and $|g_1\rangle$.
- ➡ We call Adiabatic Evolution the process represented by:

$$H_{ad}(\tau; \theta) = [1 - s(\tau; \theta)]H_0 + s(\tau; \theta)H_1. \quad (1)$$

- ➡ in which we parametrize the **scheduling** function s .

Adiabatic Evolution

- ➡ Consider two Hamiltonians H_0 and H_1 , whose ground states are $|g_0\rangle$ and $|g_1\rangle$.
- ➡ We call Adiabatic Evolution the process represented by:

$$H_{ad}(\tau; \theta) = [1 - s(\tau; \theta)]H_0 + s(\tau; \theta)H_1. \quad (1)$$

- ➡ in which we parametrize the **scheduling** function s .

```
1  # with qibo we can implement an Adiabatic Evolution via trotter formula
2  from qibo import models, hamiltonians, callbacks
3
4  # problem's parameters
5  nqubits = 1
6  h0 = hamiltonians.X(nqubits)
7  h1 = hamiltonians.Z(nqubits)
8  target_observable = h1
9
10 # we track the energy of h1 on the evolved ground state
11 energies = callbacks.Energy(target_observable)
12 evolution = models.AdiabaticEvolution(
13     h0=h0, h1=h1, s = lambda t : t, dt=0.1, callbacks = [energies])
14
15 # calculate the evolved final state at time t=final_time
16 evolved_state = evolution(final_time = final_time)
```

Machine Learning

\mathcal{M} : model;

\mathcal{O} : optimizer;

\mathcal{J} : loss function.

(x, y) : data

Quantum Computation

\mathcal{Q} : qubits;

\mathcal{S} : superposition;

\mathcal{E} : entanglement.

Machine Learning

\mathcal{M} : model;

\mathcal{O} : optimizer;

\mathcal{J} : loss function.

(x, y) : data

Quantum Computation

\mathcal{Q} : qubits;

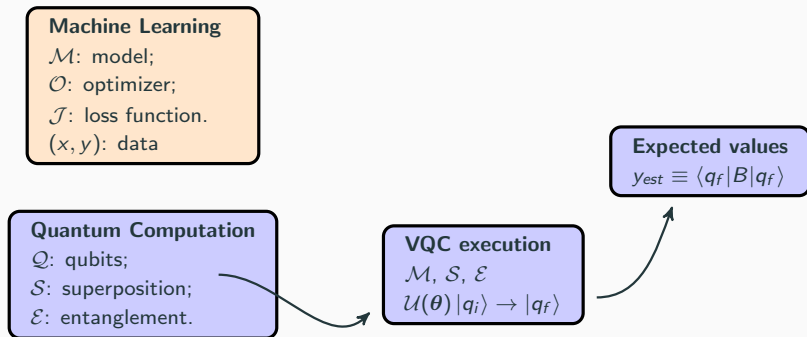
\mathcal{S} : superposition;

\mathcal{E} : entanglement.

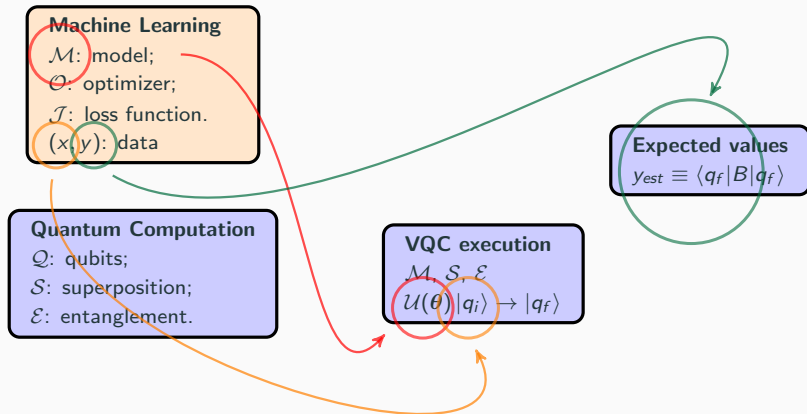
VQC execution

$\mathcal{M}, \mathcal{S}, \mathcal{E}$

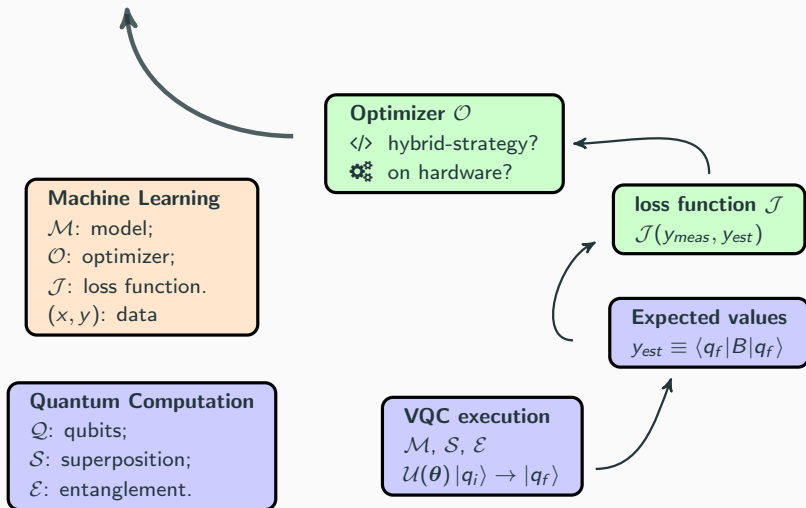
$\mathcal{U}(\theta) |q_i\rangle \rightarrow |q_f\rangle$



Quantum Machine Learning - encoding the problem



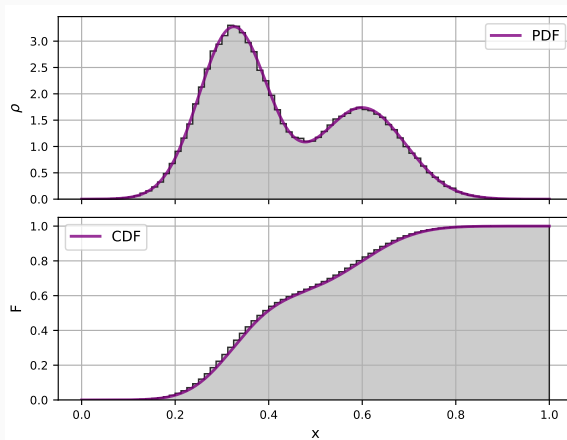
Quantum Machine Learning!



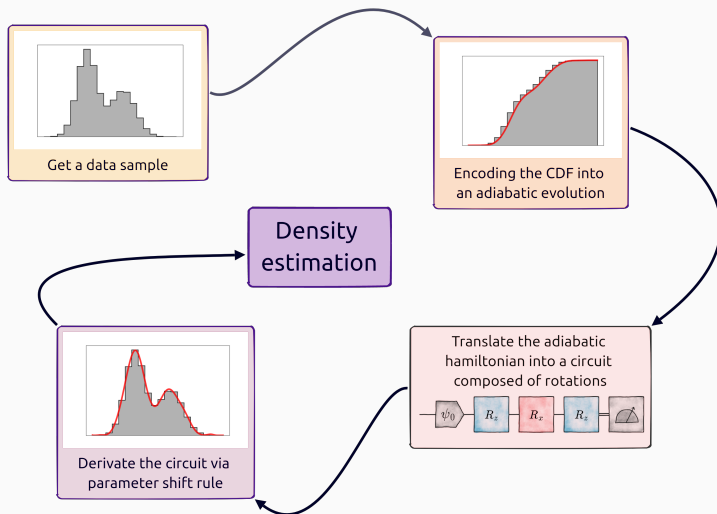
Full stack algorithms using qibo

STEP 0: the goal

- ➔ Let's consider a sample of data $\{x\}_{k=1}^{N_{\text{data}}}$.
- ➔ We can calculate the Cumulative Density Function (CDF) values $F(x)$,
- ➔ which are related to the Probability Density Function (PDF) via $\rho(x) = dF(x)/dx$.



STEP 0: the idea



- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
- 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

Training the evolution

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

Training the evolution

1. we run the evolution with random initial θ_0 into the scheduling;

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

Training the evolution

1. we run the evolution with random initial θ_0 into the scheduling;
2. we track the energy of a Pauli Z during the evolution;

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

Training the evolution

1. we run the evolution with random initial θ_0 into the scheduling;
2. we track the energy of a Pauli Z during the evolution;
3. we calculate a loss function J_{mse} :

$$J_{\text{mse}} = \sum_{k=1}^{N_{\text{sample}}} [E(\tau_k) - F(x_k)]^2;$$

- ➔ Given a sample $\{x\}$ and calculated its CDF values $F(x)$:
 - 🔗 we select two hamiltonians H_0 and H_1 such that a target observable has energy $E = 0$ and $E = 1$ respectively on H_0 and H_1 ground states;
 - 🔗 we map $(x, F) \rightarrow (\tau, E)$.

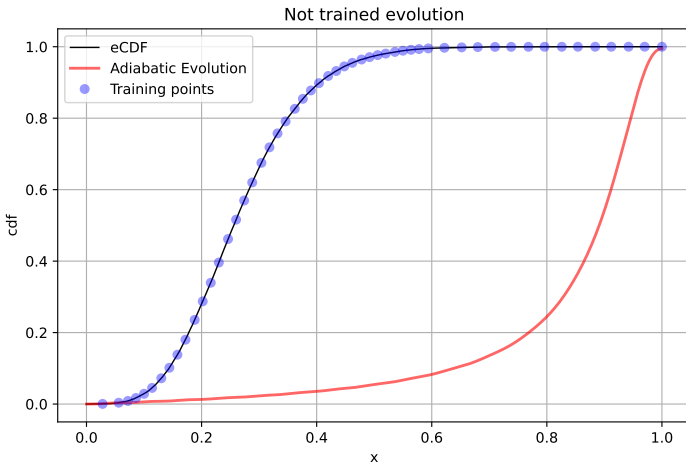
Training the evolution

1. we run the evolution with random initial θ_0 into the scheduling;
2. we track the energy of a Pauli Z during the evolution;
3. we calculate a loss function J_{mse} :

$$J_{\text{mse}} = \sum_{k=1}^{N_{\text{sample}}} [E(\tau_k) - F(x_k)]^2;$$

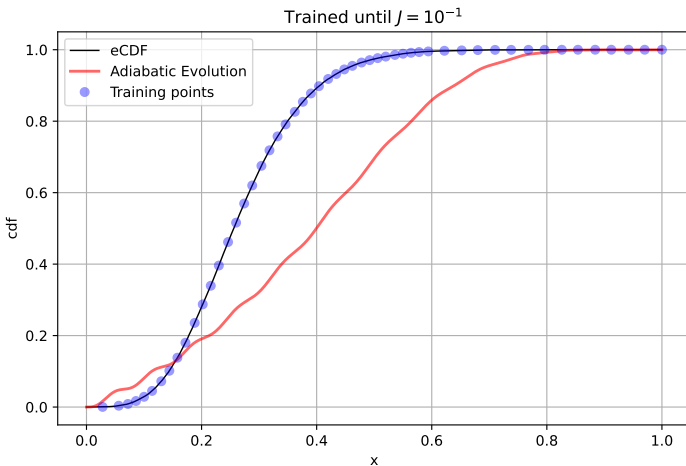
4. we choose an optimizer to find θ_{best} which minimizes J_{mse} .

➔ `nparams=20, dt=0.1, final_time=50, target_loss=None`

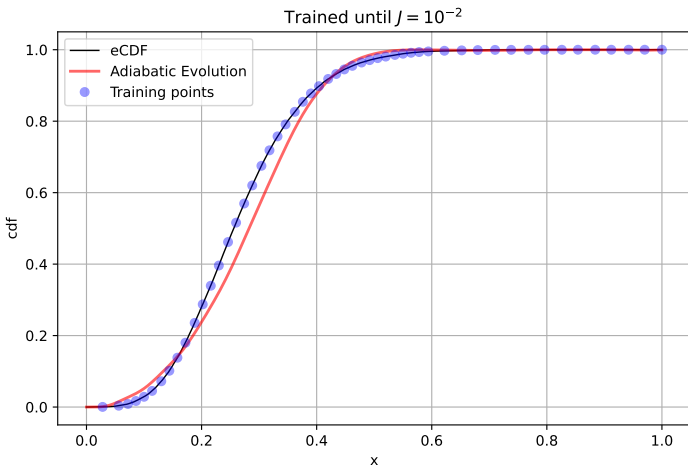


A toy example - until $J_{\text{MSE}} = 10^{-1}$

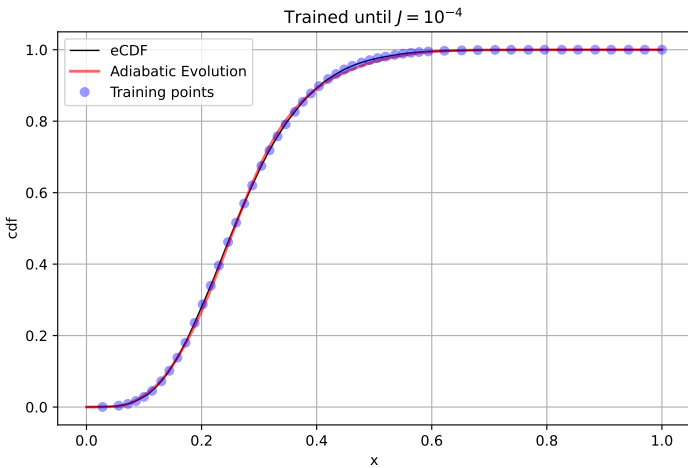
➔ `nparams=20, dt=0.1, final_time=50, target_loss=1e-1`



➔ `nparams=20, dt=0.1, final_time=50, target_loss=1e-2`



➔ `nparams=20, dt=0.1, final_time=50, target_loss=1e-4`



- ➔ Firstly, we did some calculations and approximations in order to:

➡ Firstly, we did some calculations and approximations in order to:

1. translate the Hamiltonians' sequence into a single unitary:

$$\prod_{j=1}^n e^{-iH_j dt} \rightarrow \mathcal{U}(t);$$

➡ Firstly, we did some calculations and approximations in order to:

1. translate the Hamiltonians' sequence into a single unitary:

$$\prod_{j=1}^n e^{-iH_j dt} \rightarrow \mathcal{U}(t);$$

2. translate this unitary in a sequence of rotational gates:

$$\mathcal{U}(t) = R_z(\theta_1)R_x(\theta_2)R_z(\theta_3) \quad \text{with } \theta_i \equiv \theta_i(t).$$

➡ Firstly, we did some calculations and approximations in order to:

1. translate the Hamiltonians' sequence into a single unitary:

$$\prod_{j=1}^n e^{-iH_j dt} \rightarrow \mathcal{U}(t);$$

2. translate this unitary in a sequence of rotational gates:

$$\mathcal{U}(t) = R_z(\theta_1)R_x(\theta_2)R_z(\theta_3) \quad \text{with } \theta_i \equiv \theta_i(t).$$

➡ Then, we derivate the expected values using parameter shift rule and chain rule.

➔ Firstly, we did some calculations and approximations in order to:

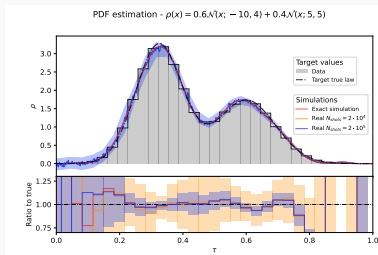
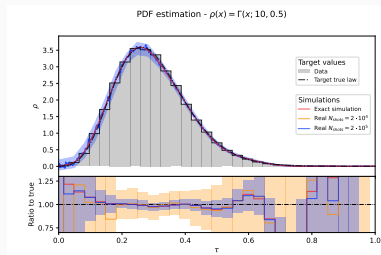
1. translate the Hamiltonians' sequence into a single unitary:

$$\prod_{j=1}^n e^{-iH_j dt} \rightarrow \mathcal{U}(t);$$

2. translate this unitary in a sequence of rotational gates:

$$\mathcal{U}(t) = R_z(\theta_1)R_x(\theta_2)R_z(\theta_3) \quad \text{with } \theta_i \equiv \theta_i(t).$$

➔ Then, we derivate the expected values using parameter shift rule and chain rule.



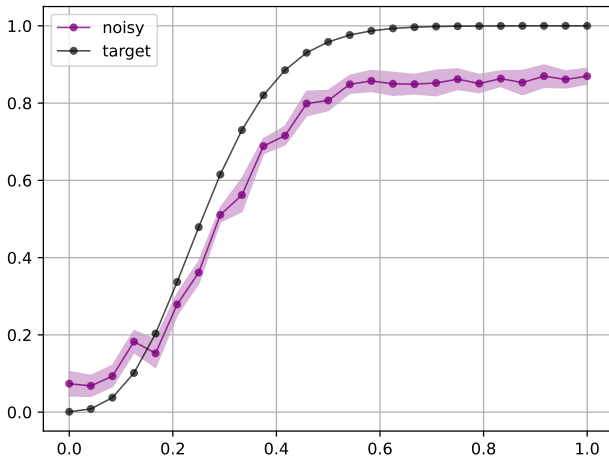


Figure 1: $N_{\text{runs}} = 10$ evaluations of CDF predictions for $N_{\text{data}} = 25$ using $N_{\text{shots}} = 1000$.

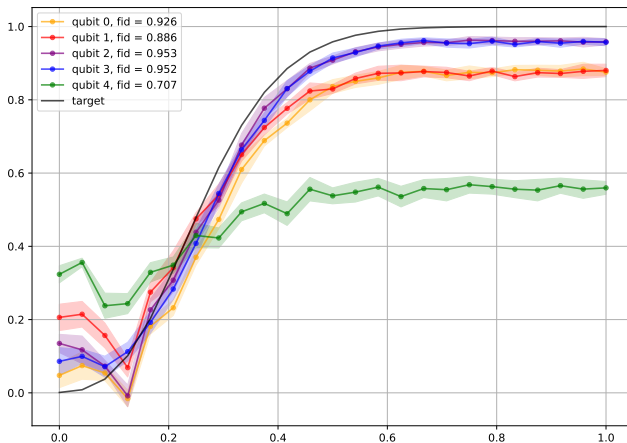


Figure 2: $N_{\text{runs}} = 10$ evaluations of CDF predictions for $N_{\text{data}} = 25$ using $N_{\text{shots}} = 1000$ and each qubit of qw5q_gold.

- ➔ These first results open several questions:

- ➔ These first results open several questions:
 - are these errors compatible with transmon SoA?

- ➔ These first results open several questions:
 - are these errors compatible with transmon SoA?
 - which relationship between my results' error and hardware ones?

- ➔ These first results open several questions:
 - are these errors compatible with transmon SoA?
 - which relationship between my results' error and hardware ones?
 - where are the execution time bottlenecks?

- ➔ These first results open several questions:
- are these errors compatible with transmon SoA?
 - which relationship between my results' error and hardware ones?
 - where are the execution time bottlenecks?
 - what if we train the entire process on the hardware?

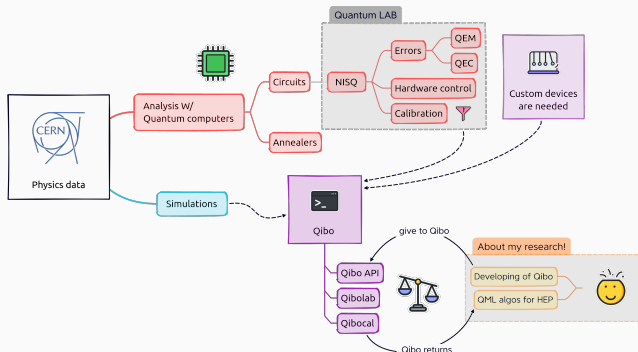
- ➡ These first results open several questions:
- are these errors compatible with transmon SoA?
 - which relationship between my results' error and hardware ones?
 - where are the execution time bottlenecks?
 - what if we train the entire process on the hardware?
 - what if we apply error mitigation on predictions?

- ➔ These first results open several questions:
- are these errors compatible with transmon SoA?
 - which relationship between my results' error and hardware ones?
 - where are the execution time bottlenecks?
 - what if we train the entire process on the hardware?
 - what if we apply error mitigation on predictions?
 - what if we apply real-time error mitigation?

The importance of the qiboteam

➡ These first results open several questions:

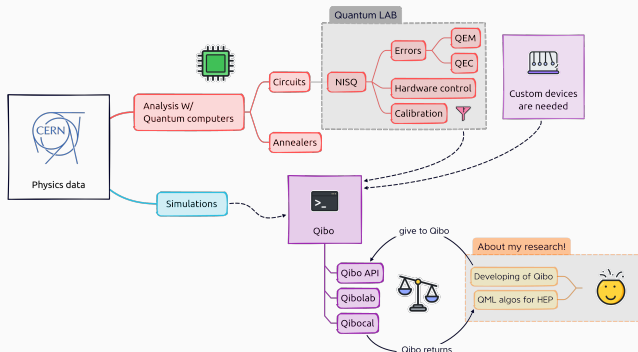
- are these errors compatible with transmon SoA?
- which relationship between my results' error and hardware ones?
- where are the execution time bottlenecks?
- what if we train the entire process on the hardware?
- what if we apply error mitigation on predictions?
- what if we apply real-time error mitigation?



The importance of the qiboteam

➡ These first results open several questions:

- are these errors compatible with transmon SoA?
- which relationship between my results' error and hardware ones?
- where are the execution time bottlenecks?
- **what if we train the entire process on the hardware?**
- what if we apply error mitigation on predictions?
- what if we apply real-time error mitigation?



- ➔ Following *Pérez-Salinas et al.* procedure, we can build a *universal quantum regressor* for approximating $y = f(x)$.

- ➔ Following *Pérez-Salinas et al.* procedure, we can build a *universal quantum regressor* for approximating $y = f(x)$.
- ➔ The model can be:

- ➔ Following *Pérez-Salinas et al.* procedure, we can build a *universal quantum regressor* for approximating $y = f(x)$.
- ➔ The model can be:

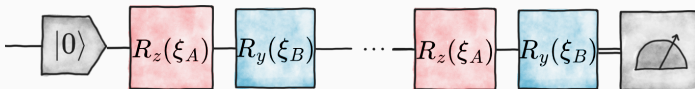


Figure 3: Here $\xi_A = \theta_1 x + \theta_2$ and $\xi_B = \theta_3 x + \theta_4$.

➔ Following *Pérez-Salinas et al.* procedure, we can build a *universal quantum regressor* for approximating $y = f(x)$.

➔ The model can be:



Figure 3: Here $\xi_A = \theta_1 x + \theta_2$ and $\xi_B = \theta_3 x + \theta_4$.

➔ and then use some $E[\hat{O}]$ as predictor:

$$y_{pred} = \langle 0 | C^\dagger(x; \theta) \hat{O} C(x; \theta) | 0 \rangle. \quad (2)$$

➔ Following *Pérez-Salinas et al.* procedure, we can build a *universal quantum regressor* for approximating $y = f(x)$.

➔ The model can be:

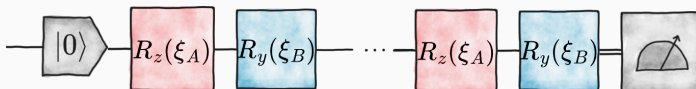


Figure 3: Here $\xi_A = \theta_1 x + \theta_2$ and $\xi_B = \theta_3 x + \theta_4$.

➔ and then use some $E[\hat{O}]$ as predictor:

$$y_{pred} = \langle 0 | C^\dagger(x; \theta) \hat{O} C(x; \theta) | 0 \rangle. \quad (2)$$

➔ Using the parameter-shift rule, we can perform a gradient descent on `tii1q_b1`.

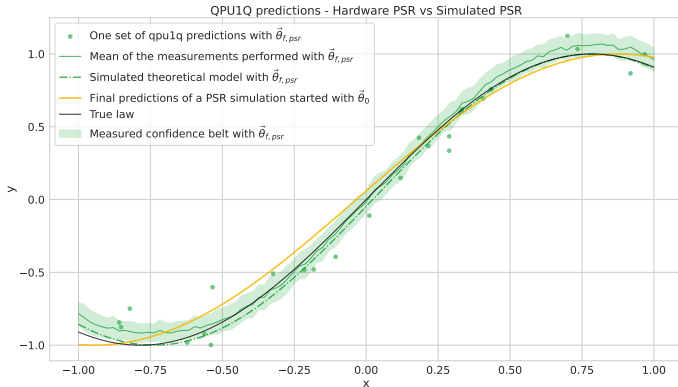


Figure 4: Batch Gradient Descent on the hardware, with gradients evaluated via Parameter-Shift Rule. We take 100 points $\{x_j\}$ in the range $[-1, 1]$ and we make 100 predictions for each x_j . Mean and standard deviation are used for determining the estimations and the confidence belt.

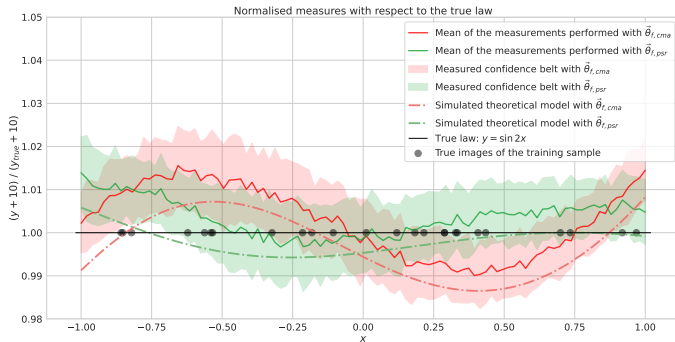


Figure 5: Normalised results of the SGD (green line) compared with true law and a genetic optimizer (red line).

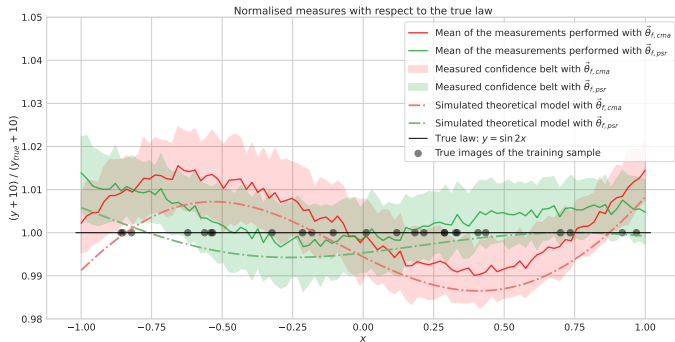


Figure 5: Normalised results of the SGD (green line) compared with true law and a genetic optimizer (red line).

👍 the full-stack framework works! comparable with a genetic algorithm;

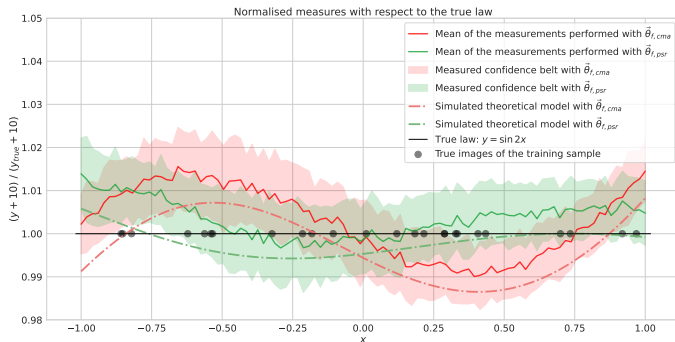


Figure 5: Normalised results of the SGD (green line) compared with true law and a genetic optimizer (red line).

- 👍 the full-stack framework works! comparable with a genetic algorithm;
- 👎 we can tackle only easy problems: it is slow;

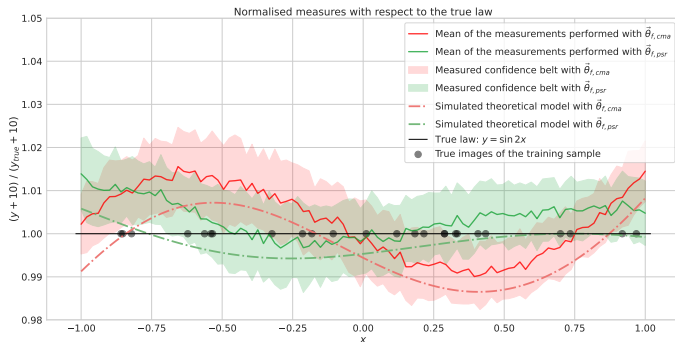


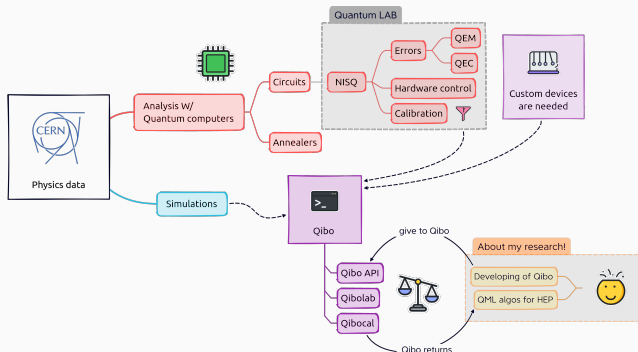
Figure 5: Normalised results of the SGD (green line) compared with true law and a genetic optimizer (red line).

- 👍 the full-stack framework works! comparable with a genetic algorithm;
- 👎 we can tackle only easy problems: it is slow;
- 😞 no mitigation: have been the errors absorbed into the optimization?

The importance of the qiboteam

➡ These first results open several questions:

- are these errors compatible with transmon SoA?
- which relationship between my results' error and hardware ones?
- where are the execution time bottlenecks?
- what if we train the entire process on the hardware?
- **what if we apply error mitigation on predictions?**
- what if we apply real-time error mitigation?



Mitigated prediction of the PDF

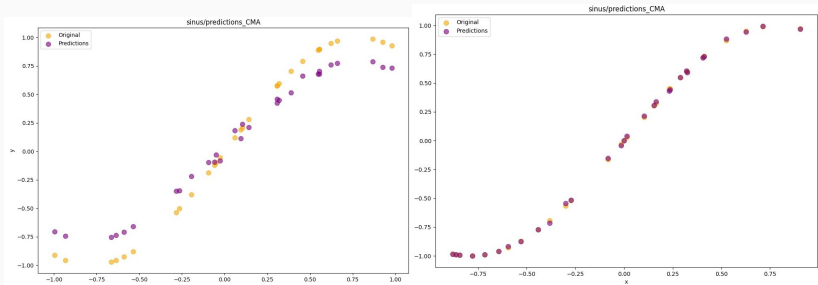


Figure 6: Predictions of a sinus function performed on `tii1q_b1` without and with CDR mitigation.

➡ These results open to new questions!

Mitigated prediction of the PDF

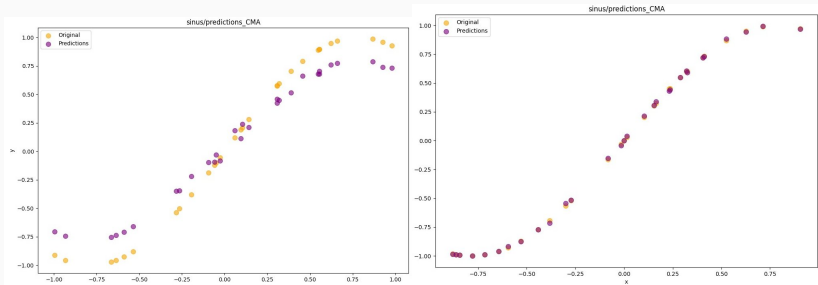


Figure 6: Predictions of a sinus function performed on tii1q_b1 without and with CDR mitigation.

➡ These results open to new questions!

- can QEM help the training with a faster convergence?

Mitigated prediction of the PDF

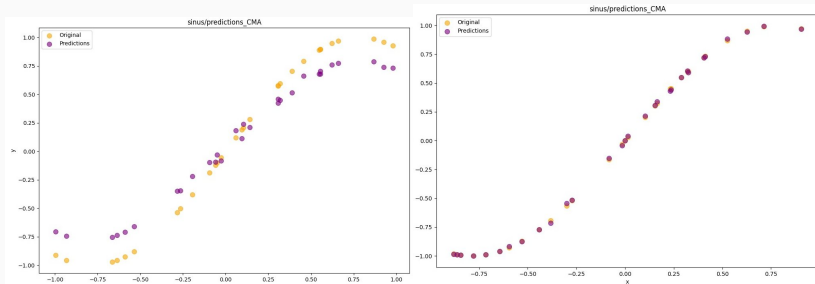


Figure 6: Predictions of a sinus function performed on tii1q_b1 without and with CDR mitigation.

➡ These results open to new questions!

- can QEM help the training with a faster convergence?
- what's the right balance between error absorption and QEM?

Mitigated prediction of the PDF

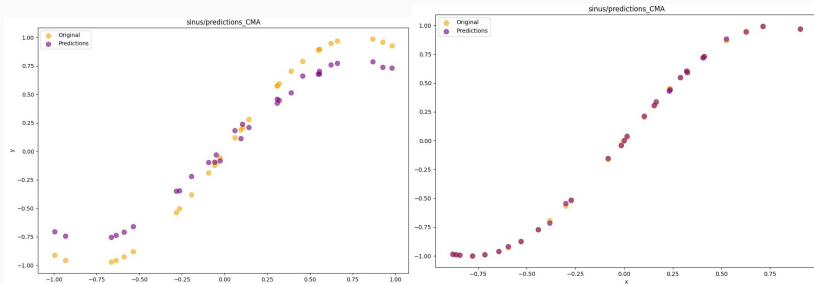


Figure 6: Predictions of a sinus function performed on tii1q_b1 without and with CDR mitigation.

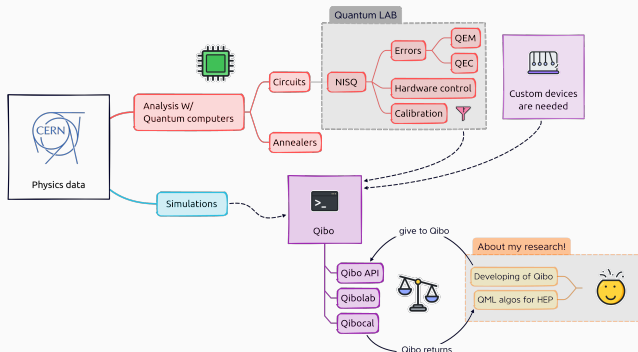
➡ These results open to new questions!

- can QEM help the training with a faster convergence?
- what's the right balance between error absorption and QEM?
- can we train the unmitigated hardware to be self-resistant?

The importance of the qiboteam

➡ These first results open several questions:

- are these errors compatible with transmon SoA?
- which relationship between my results' error and hardware ones?
- where are the execution time bottlenecks?
- what if we train the entire process on the hardware?
- what if we apply error mitigation on predictions?
- **what if we apply real-time error mitigation?**



- ➔ We want to reproduce the u quark PDF fit of *Pérez-Salinas et al.*

¹We used Zero Noise Extrapolation (ZNE) and Clifford Data Regression (CDR).

- ➔ We want to reproduce the u quark PDF fit of *Pérez-Salinas et al.*
- ➔ We apply error mitigation techniques¹ during a QML training!

¹We used Zero Noise Extrapolation (ZNE) and Clifford Data Regression (CDR).

- ➔ We want to reproduce the u quark PDF fit of *Pérez-Salinas et al.*
- ➔ We apply error mitigation techniques¹ during a QML training!

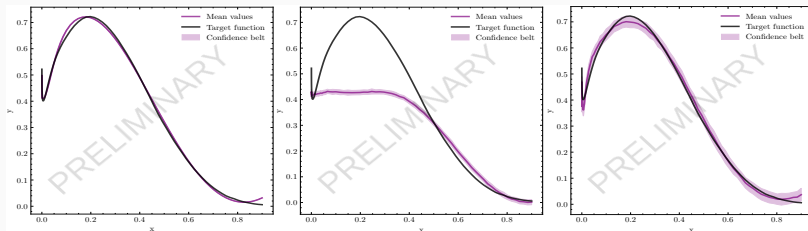


Figure 7: PDF fit performed with different levels of noisy simulation. From left to right, exact simulation, noisy simulation, noisy simulation applying error mitigation to the predictions.

¹We used Zero Noise Extrapolation (ZNE) and Clifford Data Regression (CDR).

- ➔ We want to reproduce the u quark PDF fit of *Pérez-Salinas et al.*
- ➔ We apply error mitigation techniques¹ during a QML training!

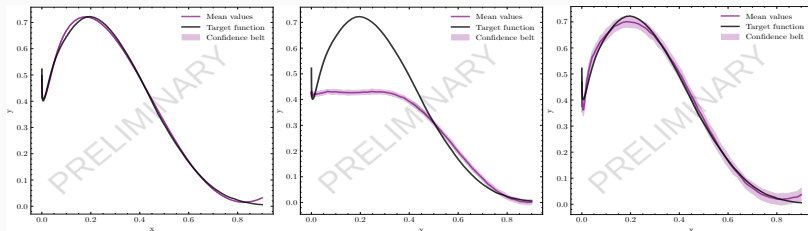


Figure 7: PDF fit performed with different levels of noisy simulation. From left to right, exact simulation, noisy simulation, noisy simulation applying error mitigation to the predictions.

- ➔ **But on the hardware?**

¹We used Zero Noise Extrapolation (ZNE) and Clifford Data Regression (CDR).

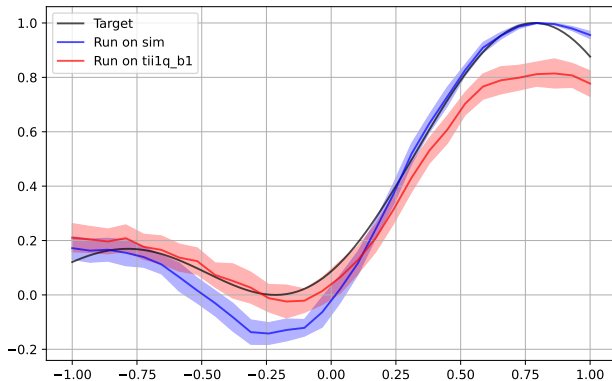


Figure 8: Full training is here performed on tii1q_b1. We then used the θ_{best} to make statistics on 30 points and 30 runs with simulation (blue line) and on the device (red line).

Conclusions

- ➔ The qibo environment is perfect for this kind of research, because

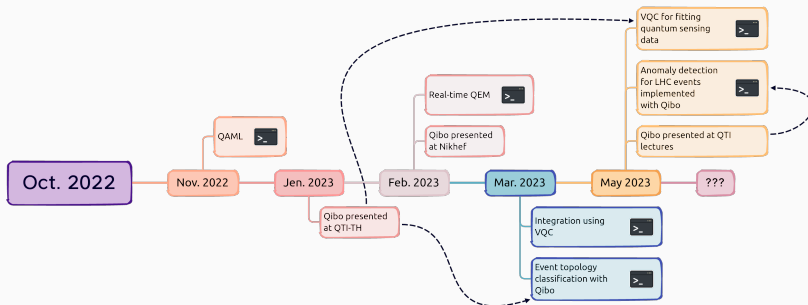
- ➡ The qibo environment is perfect for this kind of research, because
 - ⚙️ we have resources deputed to the research;

- ➔ The qibo environment is perfect for this kind of research, because
 - ⚙️ we have resources deputed to the research;
 - 👤 we are a fully connected research network.

- ➔ The qibo environment is perfect for this kind of research, because
 - ⚙️ we have resources deputed to the research;
 - 👤 we are a fully connected research network.
- ➔ Moreover, qibo is growing as international environment:

qibo as research network!

- ➔ The qibo environment is perfect for this kind of research, because
 - ⚙️ we have resources deputed to the research;
 - 👤 we are a fully connected research network.
- ➔ Moreover, qibo is growing as international environment:



Thanks!

